



Software Consulting



DevOps Projects



Corporate Training

“**UpGrade Your Business with
BEST BRAINS**”



Docker



Kubernetes



Rajesh G

Master Trainer & CTO,
Brain Upgrade Academy,
A division of Unisuraksha Tracking Systems Pvt Ltd

Brief about me



Rajesh Ghaware

Author | TOGAF EA | CKA | Trainer (Cloud - Kubernetes & Docker) | Entrepreneur | Technology Advisor

Bangalore Urban, Karnataka, India · [500+ connections](#)

My Certifications

- # CKA: Certified Kubernetes Administrator (2020)
- # TOGAF Certified Enterprise Architect (2015)
- # Spring – Core (3.2) Certified Professional (96% score – Aug 2013)
- # OMG UML (FL) – 2010
- # IBM OOAD-UML / Brainbench UML Certification (<http://www.brainbench.com/transcript.jsp?pid=3730016>) - 2005
- # Brainbench JAVA-EJB 2.0 Certification – 2001

My Interest Areas:

- # Kubernetes, Docker, AWS, IoT, Enterprise Architecture, Spring Framework, Java, UML

Some of my articles on Dzone.com as reference :

- <https://dzone.com/articles/microservices-with-observability-on-kubernetes>
- <https://dzone.com/articles/scalable-jenkins-on-kubernetes-cluster-amp-pipeline>
- <https://dzone.com/articles/expose-your-app-to-the-internet-using-ingress-cont>

Misc:

- # ICC & IE and IEMS published my research work on Production Planning and Control (2001)
- # My articles in Open source magazine - <http://opensourceforu.com/author/rajesh-ghaware>

Experience (20+ years in IT - Software Engineering)



UniGPS Solutions

3 yrs 9 mos

Founder & CTO

Dec 2018 – Present · 2 yrs 5 mos
Bengaluru Area, India

Managing Partner, CTO

Aug 2017 – Present · 3 yrs 9 mos
Bangalore

<https://unigps.in>



VP Technology

JPMorgan Chase & Co.
Jun 2015 – Aug 2017 · 2 yrs 3 mos
Bangalore Urban, Karnataka, India



Lead Solution Architect / Delivery Manager

Deutsche Bank Group
Dec 2009 – Jun 2015 · 5 yrs 7 mos
Part of OTC Derivatives since 2010



Technical Project Manager / Application Architect

Headstrong India Pvt Ltd
Jun 2006 – Dec 2009 · 3 yrs 7 mos

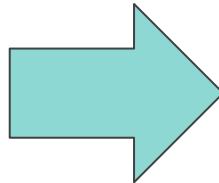
Kubernetes





Training Objectives

At the end of training,
participants should be able to



- Know Kubernetes and Be a Helmsman
- Create and run PODs
- Bundle applications & Deploy
- Service apps using Load Balancers
- Troubleshoot



Kubernetes Core

- Why Kubernetes
- Architecture
- Core Components
- API Primitives
- Kubectl
- Demo
- Practicals



What is / why Kubernetes

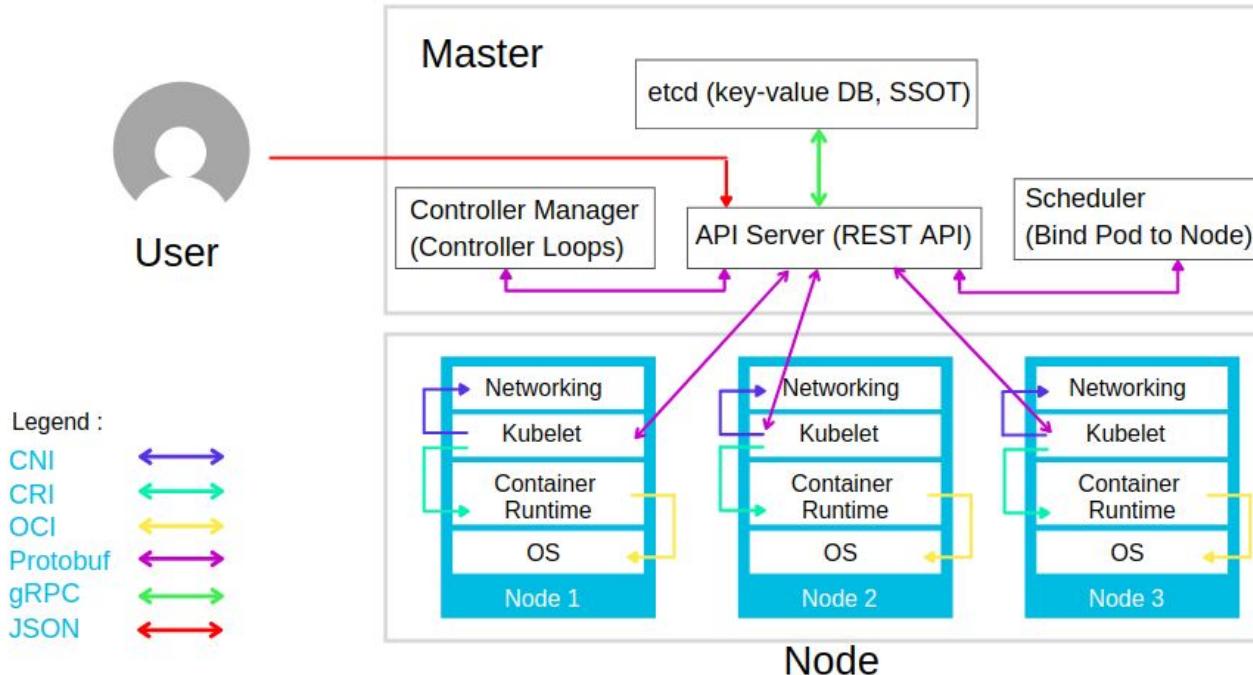
Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation.

Why?

- Service Discovery & Load Balancing
- Storage Orchestration
- Automated rollouts & rollbacks
- Automatic bin packing
- Self-healing
- Secret and configuration Management

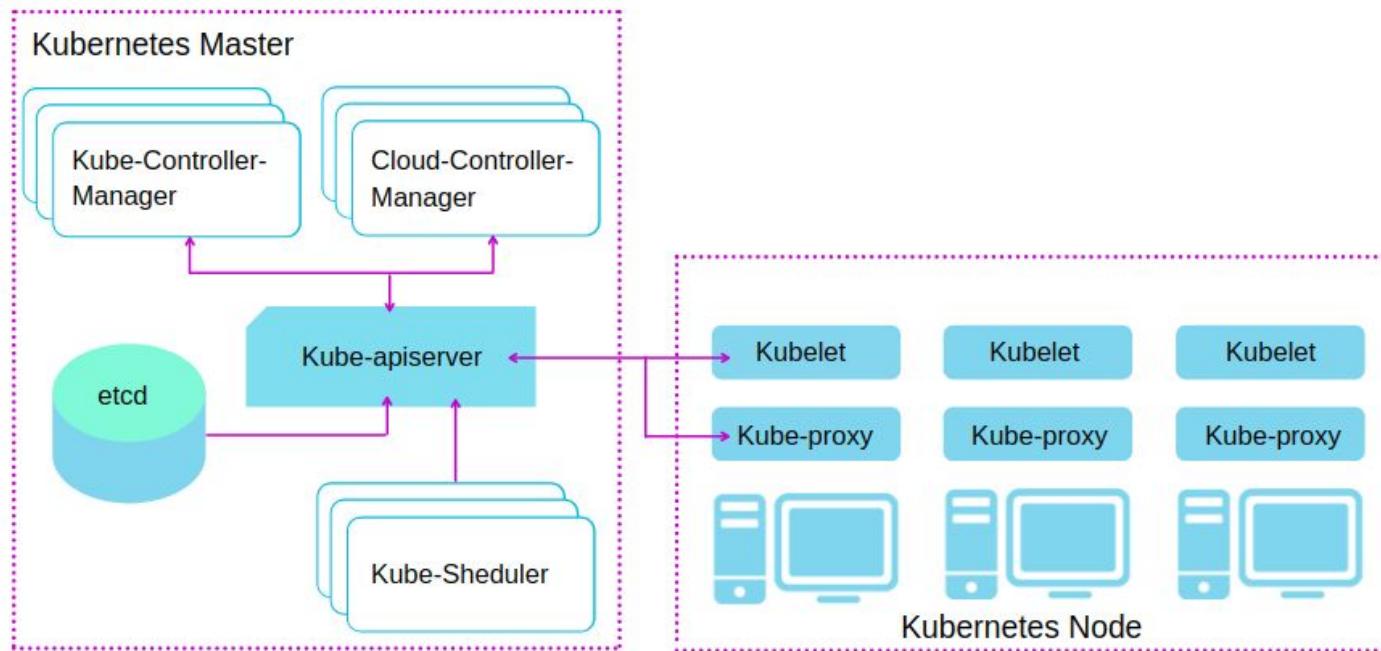


Architecture Overview





Architecture Overview





Master Components - ETCD

- Distributed reliable key-value store that is simple, secure & fast
- Uses RAFT based consensus algorithm to work in distributed environment
- Key value store distributed database
- Runs on port 2379



Master Components - API Server

- The central management entity
- Only component that connects to ETCD
- Designed for horizontal scaling

Connectivity:

- External: kubectl
- Internal: kubelet
- Persistent Storage: ETCD



Master Components - Scheduler

Schedules pods on appropriate Node(s)

Watches for newly created PODs that have no nodes assigned

Decision Parameters:

- Resource requirements (memory, cpu, disk type say SSD)
- Hardware, Software, Policy requirements
- Affinity, Anti-affinity
- Data locality



Master Components - Kube Controller

- Replication Controller
 - Responsible for maintaining the correct number of pods for every replication controller object in the system
- Endpoints Controller
 - Populates the Endpoints object (that is, joins Services & Pods)
- Service Account & token Controller
 - Create default accounts and API access tokens for new namespaces



Master Components - Cloud Controller

- Node Controller
 - Updates Node objects when new servers are created in the cloud, including labels, hostname, network addresses, verify health etc
- Route Controller
 - For setting up routes in the underlying cloud infrastructure so that containers on different nodes can communicate with each other
- Service Controller
 - For creating, updating and deleting cloud provider load balancers
- Volume Controller
 - For creating, attaching, and mounting volumes, and interacting with the cloud provider to orchestrate volumes



Node Components

- KUBE-PROXY
 - Network proxy that runs on every node in cluster
 - Maintains network rules on nodes
 - Uses OS packet filtering layer else forwards traffic itself
- KUBELET
 - Runs on every node
 - Ensures containers are running & healthy in PODs
 - Doesn't manage container not created by K8S



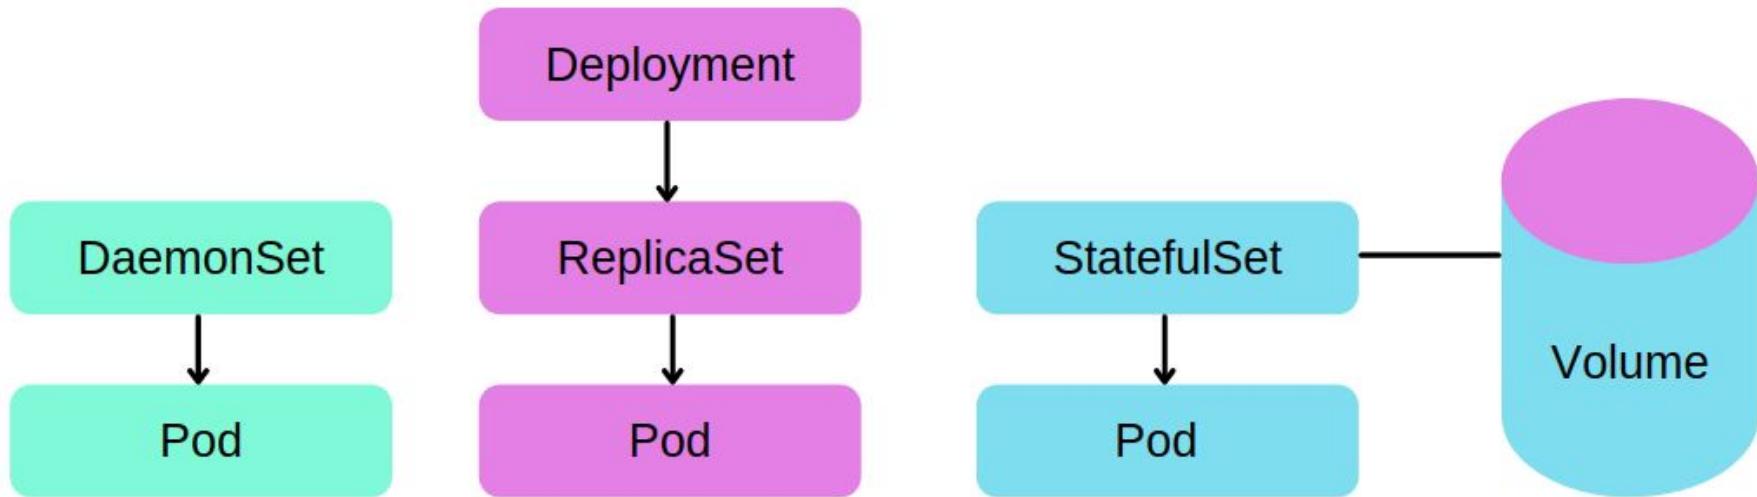
Kubectl

Command line tool to control kubernetes cluster

- Imperative commands to manage objects (basic & intermediate)
- Deploy commands
- Cluster Management commands
- Troubleshooting and Debugging
- Advanced, Settings and Other



Objects





kubectl - commands

- kubectl get pods
- kubectl describe pod hello-world
- kubectl describe pod/nginx
- kubectl delete pod nginx
- kubectl cluster-info
- kubectl get pods -o yaml
- kubectl get services -o json
- kubectl get pods --sort-by=.metadata.name
- kubectl get rs,deployments,service
- kubectl describe pods
- kubectl get pod/<pod-name> svc/<svc-name>
- kubectl get pod -l name=<label-name>
- kubectl delete pods --all
- `kubectl get nodes -o json | jq '.items[] | {name:.metadata.name, cap:.status.capacity}'`
- `kubectl get nodes -o yaml | egrep '\sname:|cpu:|memory:'`
- kubectl get all



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



POD

- Overview
- Lifecycle
- Init Containers
- Topology Spread



POD - Overview

- ❖ Smallest deployable unit
- ❖ Supports multiple cooperating processes (containers) that form cohesive unit of service
- ❖ Ephemeral Entity

Encapsulates

- application container(s)
- Storage resources
- Unique network IP

Shared Resources:

- Networking
- Storage



Example

```
training > kubernetes > lab > pod > ! nginx-pod.yaml
```

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    labels:
5      run: nginx
6    name: nginx
7  spec:
8    containers:
9      - image: nginx
10     name: nginx
11   dnsPolicy: ClusterFirst
12   restartPolicy: Always
```

```
training > kubernetes > lab > kubernetes-core > ! 01-monolith-pod.yaml
```

```
1  # Monolith containing all app components
2  # including database
3  apiVersion: v1
4  kind: Pod
5  metadata:
6    labels:
7      app: monolith
8      name: monolith
9  spec:
10    containers:
11      - image: brainupgrade/weather:monolith
12        name: monolith
13    restartPolicy: Always
```



Example

```
training > kubernetes > lab > pod > ! pod-busybox.yaml
 1  apiVersion: v1
 2  kind: Pod
 3  metadata:
 4    labels:
 5      run: pod-busybox
 6      name: pod-busybox
 7  spec:
 8    containers:
 9      - command:
10        - sh
11        - -c
12        - echo App is running! && sleep 30
13      image: busybox
14      name: pod-busybox
15    restartPolicy: Never
```



POD - Lifecycle ...

- Lifecycle Phases
 - Pending (waiting to be scheduled, image downloading)
 - Running (all containers started and ready to serve)
 - Succeeded (all containers exited with success)
 - Failed (all containers exited but at least one with failure)
 - Unknown (unable to fetch status as node is unreachable)
- Container States
 - Waiting, Running, Terminated
- Restart Policy (**Always**, Never, OnFailure)
- Probes
 - Startup, Readiness, Liveness
- Lifecycle hooks



Phase - Pending

```
training > kubernetes > lab > pod > ! lifecycle-phase-pending.yaml
 1  apiVersion: v1
 2  kind: Pod
 3  metadata:
 4    labels:
 5      run: nginx
 6      name: nginx
 7  spec:
 8    containers:
 9      - image: nginx
10        name: nginx
11        resources:
12          requests:
13            cpu: "6000m"
14            memory: "100Gi"
15        dnsPolicy: ClusterFirst
16        restartPolicy: Never
```



Phase - Running

```
training > kubernetes > lab > pod > ! lifecycle-phase-running.yaml
 1  apiVersion: v1
 2  kind: Pod
 3  metadata:
 4    labels:
 5      run: busybox
 6    name: busybox
 7  spec:
 8    containers:
 9      - command:
10        - ping
11        - google.com
12      image: busybox
13      name: busybox
14    dnsPolicy: ClusterFirst
15    restartPolicy: Always
```

Phase - Succeeded

```
training > kubernetes > lab > pod > ! lifecycle-phase-succeeded.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    labels:
5      run: busybox
6      name: busybox
7  spec:
8    containers:
9      - image: busybox
10     name: busybox
11   dnsPolicy: ClusterFirst
12   restartPolicy: Never
```



Probes

- Types
 - Startup
 - Readiness
 - Liveness
- Methods
 - Http
 - Tcp
 - Command
- Settings
 - initialDelaySeconds
 - periodSeconds
 - timeoutSeconds
 - successThreshold
 - failureThreshold
- Http
 - Host
 - Scheme
 - Path
 - Port
 - Headers



Probe - Liveness - Exec

```
training > kubernetes > lab > pod > ! probe-liveness-exec.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: probe-liveness-exec
5  spec:
6    containers:
7      - name: probe-liveness-exec
8        image: k8s.gcr.io/busybox
9        args:
10       - /bin/sh
11       - -c
12       - touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600
13     livenessProbe:
14       exec:
15         command:
16           - cat
17           - /tmp/healthy
18       initialDelaySeconds: 5
19       periodSeconds: 5
```



Probe - Liveness - http

```
training > kubernetes > lab > pod > ! probe-liveness-http.yaml
 1  apiVersion: v1
 2  kind: Pod
 3  metadata:
 4    name: probe-liveness-http
 5  spec:
 6    containers:
 7      - name: probe-liveness-http
 8        image: k8s.gcr.io/liveness
 9        args:
10          - /server
11        livenessProbe:
12          httpGet:
13            path: /healthz
14            port: 8080
15            httpHeaders:
16              - name: Custom-Header
17                value: Awesome
18            initialDelaySeconds: 3
19            periodSeconds: 3
```

Ref <https://github.com/kubernetes/kubernetes/blob/master/test/images/agnhost/liveness/server.go>



Probe - Liveness - readiness - tcp

```
training > kubernetes > lab > pod > ! probe-liveness-readiness-tcp.yaml
 1  apiVersion: v1
 2  kind: Pod
 3  metadata:
 4    name: probe-liveness-readiness-tcp
 5  spec:
 6    containers:
 7      - name: probe-liveness-readiness-tcp
 8        image: k8s.gcr.io/goproxy:0.1
 9        ports:
10          - containerPort: 8080
11        readinessProbe:
12          tcpSocket:
13            port: 8080
14          initialDelaySeconds: 5
15          periodSeconds: 10
16        livenessProbe:
17          tcpSocket:
18            port: 8080
19          initialDelaySeconds: 15
20          periodSeconds: 20
```



Probe - Liveness - startup - http

```
training > kubernetes > lab > pod > ! probe-liveness-startup-http.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: probe-liveness-startup-http
5  spec:
6    containers:
7      - name: probe-liveness-startup-http
8        image: k8s.gcr.io/liveness
9        args:
10          - /server
11        livenessProbe:
12          httpGet:
13            path: /healthz
14            port: 8080
15            failureThreshold: 1
16            periodSeconds: 10
17        startupProbe:
18          httpGet:
19            path: /healthz
20            port: 8080
21            failureThreshold: 30
22            periodSeconds: 10
```



POD Init Containers

- Always run to completion
- Must complete successfully before next one
- Readiness probes not supported
- Run(s) before application containers

Examples:

- Custom code / utilities to run before app containers
- Block / delay app container startup
- App container image building can be separate



Example

```
training > kubernetes > lab > pod > ! init-containers.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: init-containers
5  spec:
6    containers:
7      - name: main-container
8        image: busybox:1.28
9        command: ['sh', '-c', 'echo The app is running! && sleep 3600']
10   initContainers:
11     - name: init-service
12       image: busybox:1.28
13       command:
14         - sh
15         - -c
16         - until nslookup myservice.$(cat /var/run/secrets/kubernetes.io/serviceaccount/namespace).svc.cluster.local; do echo waiting for myservice; sleep 2; done
17     - name: init-mydb
18       image: busybox:1.28
19       command: ['sh', '-c', "until nslookup mydb.$(cat /var/run/secrets/kubernetes.io/serviceaccount/namespace).svc.cluster.local; do echo waiting for mydb; sleep 2; done"]
```



Example

```
Training > app > weather > ./orb-apm-weather-front.yaml
1  # microservices components 1 - front
2  apiVersion: apps/v1
3  kind: Deployment
4  metadata:
5    labels:
6      app: weather
7      name: weather-front
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: weather-front
13   template:
14     metadata:
15       labels:
16         app: weather-front
17   spec:
18     volumes:
19       - name: elastic-apm-agent
20         emptyDir: {}
21     initContainers:
22       - name: elastic-java-agent
23         image: docker.elastic.co/observability/apm-agent-java:1.12.0
24         volumeMounts:
25           - mountPath: /elastic/apm/agent
26             name: elastic-apm-agent
27           command: ['cp', '-v', '/usr/agent/elastic-apm-agent.jar', '/elastic/apm/agent']
28     containers:
29       - name: brainupgrade/weather:microservices.front
30         imagePullPolicy: Always
31         name: weather-front
32         volumeMounts:
33           - mountPath: /elastic/apm/agent
34             name: elastic-apm-agent
35         env:
36           - name: spring.application.name
37             value: weather-front
38           - name: spring.datasource.url
39             value: jdbc:mysql://weather-db:3306/weather
40           - name: spring.datasource.username
41             value: weather
42           - name: spring.datasource.password
43             value: weather
44           - name: weatherServiceURL
45             value: http://weather-services.weather.svc.cluster.local
46           - name: ELASTIC_APM_SERVER_URL
47             value: "http://apm-server.elasticsearch.svc.cluster.local:8200"
48           - name: ELASTIC_APM_SERVICE_NAME
49             value: "weather-front"
50           - name: ELASTIC_APM_APPLICATION_PACKAGES
51             value: "in.brainupgrade"
52           - name: ELASTIC_APM_ENVIRONMENT
53             value: prod
54           - name: ELASTIC_APM_LOG_LEVEL
55             value: DEBUG
56           - name: JAVA_TOOL_OPTIONS
57             value: -javaagent:/elastic/apm/agent/elastic-apm-agent.jar
58         ports:
59           - containerPort: 8080
```

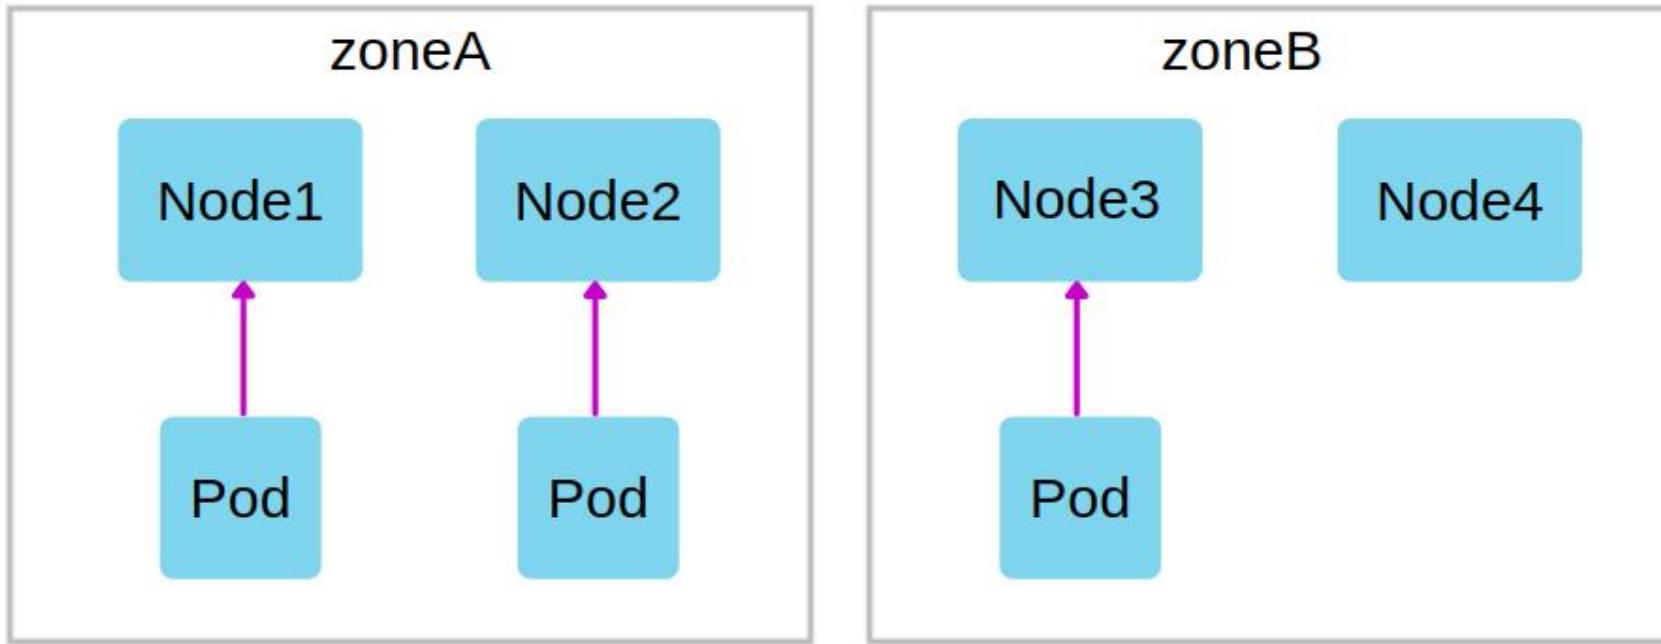


POD - Topology Spread

- Objectives
 - To control how Pods are spread across zones, nodes and other user defined topology domains
 - To achieve high availability
 - To achieve efficient resource utilization
- Spread Constraints
 - maxSkew
 - topologyKey
 - whenUnsatisfiable (DoNotSchedule / ScheduleAnyway)
 - labelSelector



POD - Topology Spread



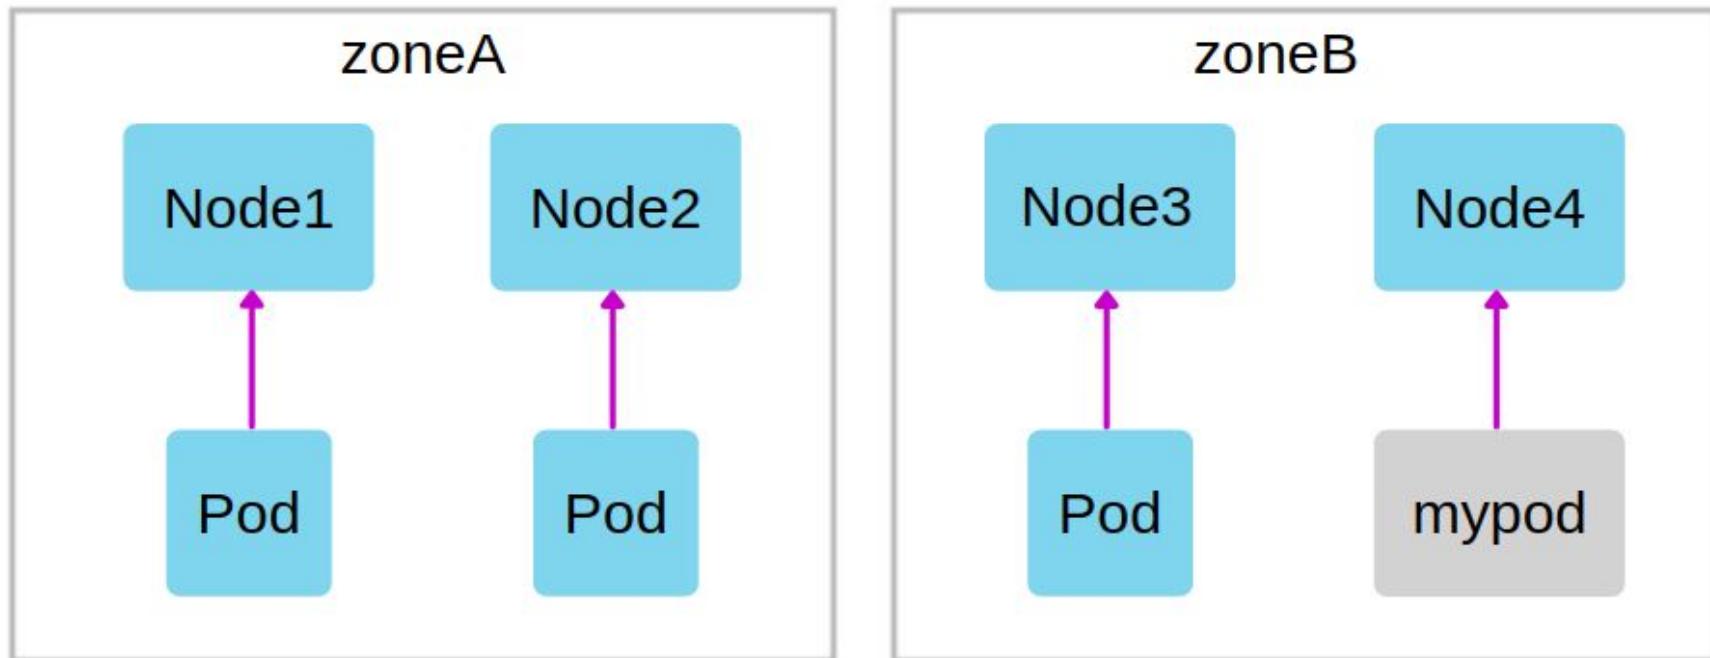


POD - Topology Spread

```
dockerk8s > kubernetes > lab > scheduling > topology > ! one-pod-per-zone.yaml >
 1  apiVersion: apps/v1
 2  kind: Deployment
 3  metadata:
 4    name: one-pod-per-zone
 5  spec:
 6    replicas: 1
 7    selector:
 8      matchLabels:
 9        app: one-pod-per-zone
10    template:
11      metadata:
12        labels:
13          app: one-pod-per-zone
14    spec:
15      nodeSelector:
16        kubernetes.io/role: node
17      topologySpreadConstraints:
18        - maxSkew: 1
19          topologyKey: topology.kubernetes.io/zone
20          whenUnsatisfiable: DoNotSchedule
21          labelSelector:
22            matchLabels:
23              app: one-pod-per-zone
24          containers:
25            - image: gcr.io/google-samples/hello-app:1.0
26              name: hello-app
```

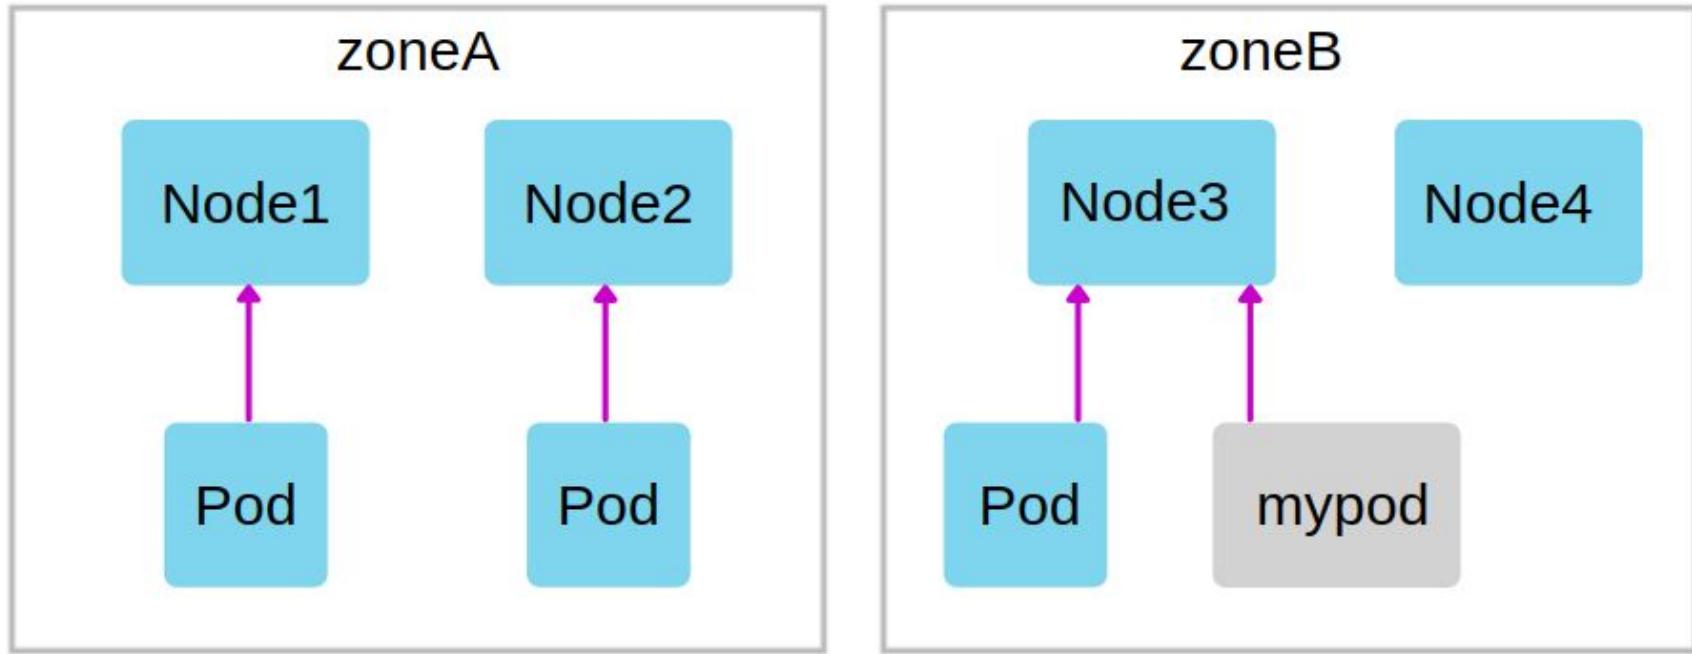


POD - Topology Spread





POD - Topology Spread





Topology - Multiple Constraints

```
dockerk8s > kubernetes > lab > scheduling > topology > ! one-pod-per-node-zone.yaml
 1 apiVersion: apps/v1
 2 kind: Deployment
 3 metadata:
 4   name: one-pod-per-node-zone
 5 spec:
 6   replicas: 1
 7   selector:
 8     matchLabels:
 9       app: one-pod-per-node-zone
10   template:
11     metadata:
12       labels:
13         app: one-pod-per-node-zone
14   spec:
15     nodeSelector:
16       kubernetes.io/role: node
17     topologySpreadConstraints:
18       - maxSkew: 1
19         topologyKey: topology.kubernetes.io/zone
20         whenUnsatisfiable: DoNotSchedule
21         labelSelector:
22           matchLabels:
23             app: one-pod-per-node-zone
24       - maxSkew: 1
25         topologyKey: kubernetes.io/hostname
26         whenUnsatisfiable: DoNotSchedule
27         labelSelector:
28           matchLabels:
29             app: one-pod-per-node-zone
30       containers:
31         - image: gcr.io/google-samples/hello-app:1.0
32           name: hello-app
```



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Pod - Advanced

- Patterns
 - Sidecar
 - Adapter
 - Ambassador
- Jobs
- Cron Jobs
- Labels, Selectors, Annotations
- Demo
- Practicals

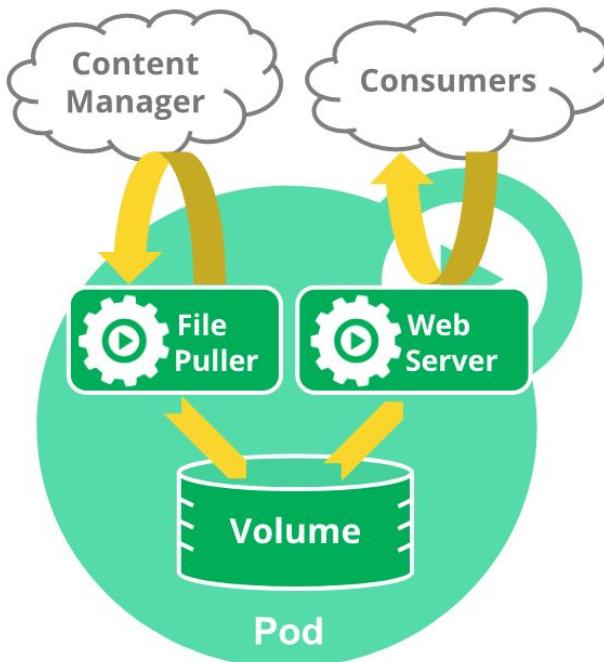


Patterns - POD

- To extend the functionality of the existing container
- To have helper process enhancing work of the existing container
- Types
 - Sidecar - To export logs, generate / pull data etc.
 - Ambassador - To proxy connection
 - Adapter - To standardise and normalize output



Pattern - Sidecar



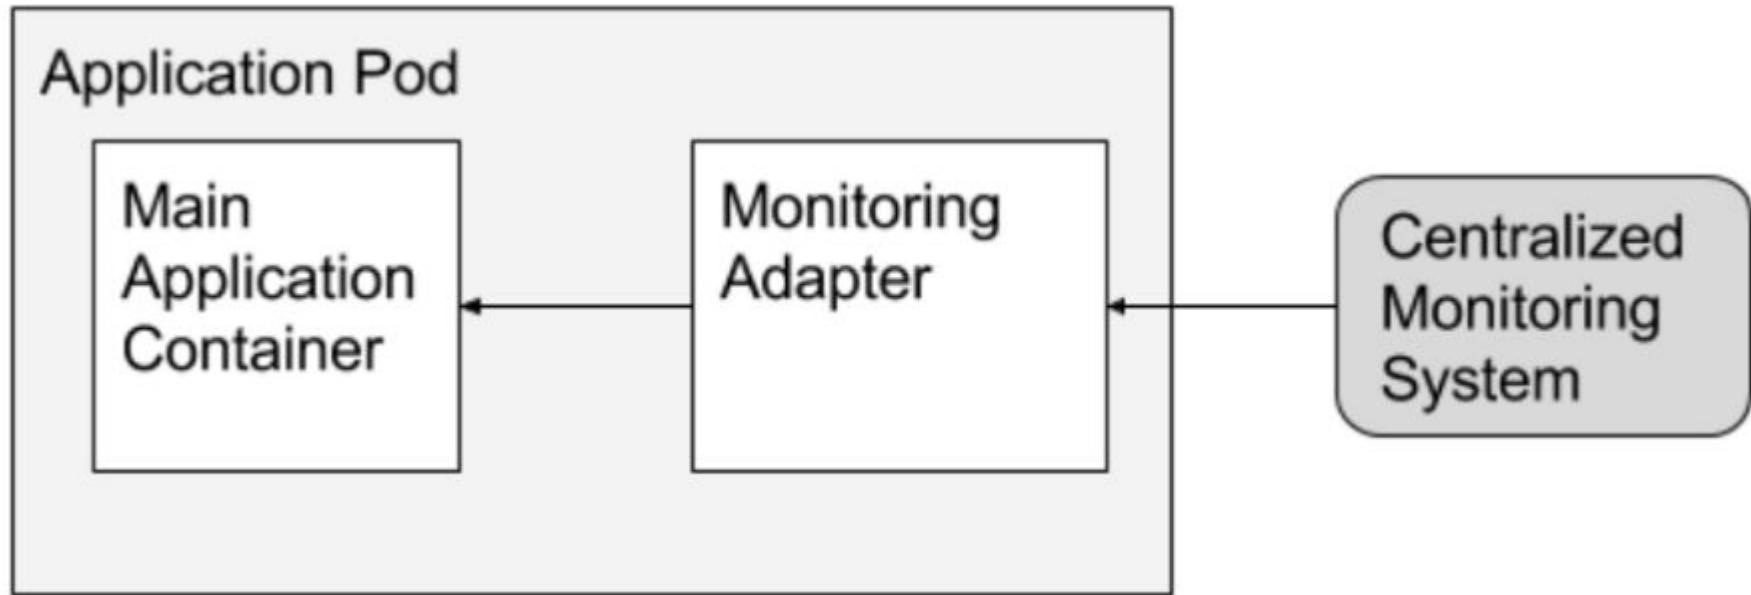


Pattern - Sidecar

```
dockerk8s > kubernetes > lab > pod-advanced > pod-multi-container > sidecar.yaml > apiVersion
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: sidecar
5  spec:
6    imagePullSecrets:
7      - name: regcred
8    volumes:
9      - name: shared-logs
10     emptyDir: {}
11   containers:
12     - name: main-container
13       image: busybox
14       command: ["/bin/sh"]
15       args: ["-c", "while true; do ping -c4 google.com >> /var/log/index-input.html; sleep 10;done"]
16       resources:
17         limits:
18           cpu: "10m"
19           memory: "10Mi"
20       volumeMounts:
21         - name: shared-logs
22           mountPath: /var/log
23     - name: sidecar-container
24       image: busybox # splunk adapter
25       resources:
26         limits:
27           cpu: "10m"
28           memory: "10Mi"
29       command: ["/bin/sh"]
30       args: ["-c", "while true; do cat /var/log/index-input.html > /tmp/index-output.html; sleep 60;done"]
31       volumeMounts:
32         - name: shared-logs
33           mountPath: /var/log
```



Pattern - Adapter



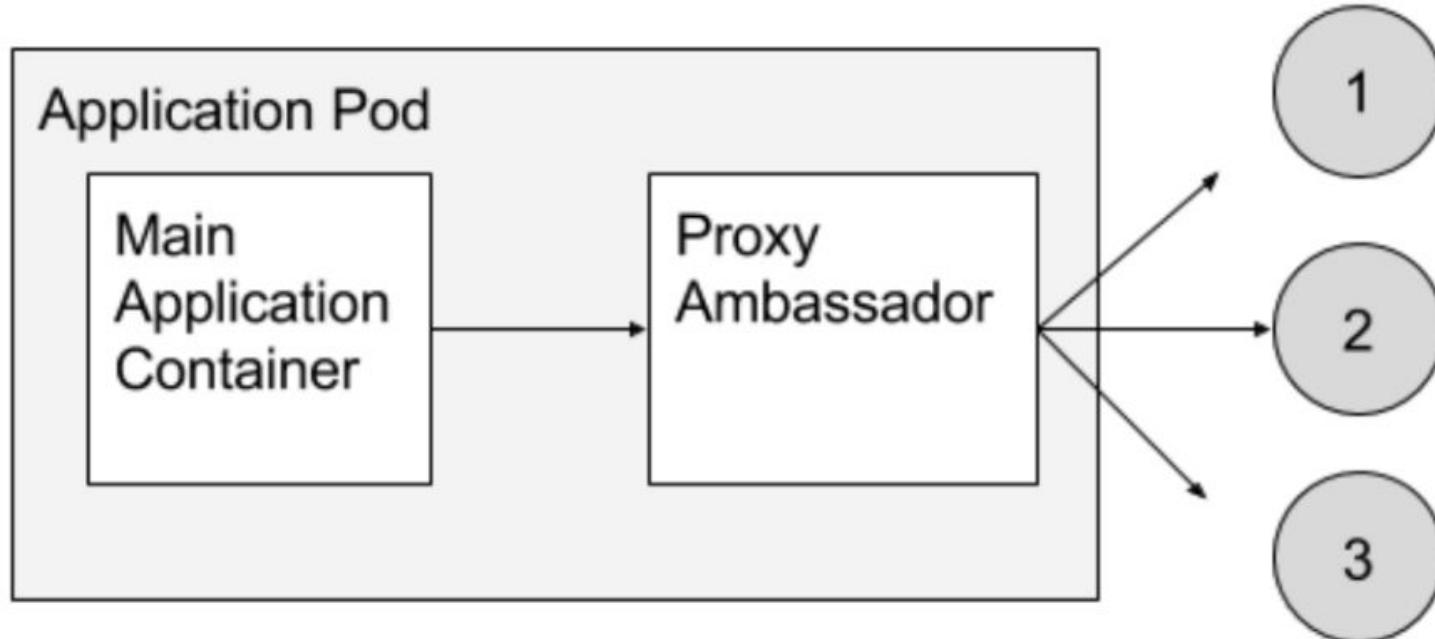


Pattern - Adapter

```
dockerk8s > kubernetes > lab > pod-advanced > pod-multi-container > adapter.yaml > [ ]-o -E '\d+\w' | head -1 >> /var/log/status.txt) && (cat /var/log/top.txt | head -3 | tail -1 | grep -o -E '\d+%' | head -1
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: adapter
5  spec:
6    imagePullSecrets:
7      - name: regcred
8    volumes:
9      - name: shared-logs
10     emptyDir: {}
11   containers:
12     - name: main-container
13       image: alpine
14       command: ["/bin/sh"]
15       args: ["-", "while true; do date > /var/log/top.txt && top -n 1 -b >> /var/log/top.txt; sleep 10;done"]
16     volumeMounts:
17       - name: shared-logs
18         mountPath: /var/log
19     - name: adapter-container
20       image: alpine # logstash message formatter
21       command: ["/bin/sh"]
22       args: ["-", "while true; do (cat /var/log/top.txt | head -1 > /var/log/status.txt) && (cat /var/log/top.txt | head -2 | tail -1 | grep
23         -o -E '\d+\w' | head -1 >> /var/log/status.txt) && (cat /var/log/top.txt | head -3 | tail -1 | grep
24         -o -E '\d+' | head -1 >> /var/log/status.txt); sleep 5; done"]
25     volumeMounts:
26       - name: shared-logs
27         mountPath: /var/log
```



Pattern - Ambassador





Pattern - Ambassador

```
training > kubernetes > lab > pod-advanced > pod-multi-container > ! ambassador.yaml
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: ambassador-nginx-config
5  data:
6    nginx.conf: |
7      worker_processes 4;
8      events { worker_connections 1024; }
9      http {
10        server {
11          listen 80;
12          location / {
13            proxy_pass https://www.brainupgrade.in;
14          }
15        }
16      }
```

```
18  apiVersion: v1
19  kind: Pod
20  metadata:
21    name: multi-pod-ambassador
22  spec:
23    imagePullSecrets:
24      - name: regcred
25    containers:
26      - name: main-app
27        image: busybox
28        imagePullPolicy: IfNotPresent
29        command: ["/bin/sh"]
30        args: ["-c","while true;do wget -O /tmp/app.txt localhost ;sleep 30;done"]
31      - name: ambassador
32        image: nginx
33        imagePullPolicy: IfNotPresent
34        ports:
35          - containerPort: 80
36        volumeMounts:
37          - name: nginx-config
38            mountPath: /etc/nginx
39        volumes:
40          - name: nginx-config
41            configMap:
42              name: ambassador-nginx-config
```



Job

- To provide reliable parallel execution of tasks
- Examples:
 - Send emails, transcode files, Scan database for a set of rows,
- Patterns
 - Parallel job
 - Fixed completion count job
 - Work queue job



Job

```
dockerk8s > kubernetes > lab > pod-advanced > ! job.yaml > ...
1  apiVersion: batch/v1
2  kind: Job
3  metadata:
4    name: migration-task
5  spec:
6    backoffLimit: 5
7    activeDeadlineSeconds: 100
8    parallelism: 10
9    template:
10      spec:
11        containers:
12          - name: cron-job
13            image: busybox
14            args:
15              - /bin/sh
16              - -c
17              - date; echo Migrating data to reporting server...;sleep 1;
18            restartPolicy: Never
```



CronJob

- Creates jobs on a repeating schedule
- Schedule times are based on kube-controller-manager
- Useful for tasks like migrating data to reporting server, sending emails, creating backups etc
- Schedule tasks at specific time (like when cluster is idle)

Key Configurations:

- startingDeadlineSeconds - Missed occurrences in last X seconds will be counted
- concurrencyPolicy
 - If Allow, then job will run at least once
 - If Forbid, will be missed if previous instance is still running



CronJob - Expression

```
# └───────── minute (0 - 59)
#   └──────── hour (0 - 23)
#     └───────── day of the month (1 - 31)
#       └──────── month (1 - 12)
#         └───────── day of the week (0 - 6) (Sunday to Saturday;
#           7 is also Sunday on some systems)
#
#
# * * * * * <command to execute>
```

Examples:

- */15 0,8,16 * * * echo running backup (every 15 minutes of 0,8 & 16th hour)
- 30 0 * * 6 /home/oracle/scripts/export_dump.sh (last day of week at 00:30)
- 1 0 * * * printf "" > /var/log/apache/error_log (everyday at 00:01)



CronJob - Expression

```
training > kubernetes > lab > pod-advanced > ! cron-job.yaml
1  apiVersion: batch/v1beta1
2  kind: CronJob
3  metadata:
4    name: cron-job
5  spec:
6    schedule: "*/10 * * * *"
7    startingDeadlineSeconds: 60
8    concurrencyPolicy: "Allow"
9    jobTemplate:
10      spec:
11        template:
12          spec:
13            imagePullSecrets:
14              - name: regcred
15            containers:
16              - name: cron-job
17                image: busybox
18                args:
19                  - /bin/sh
20                  - -c
21                  - date; echo Migrating data to reporting server...;sleep 60;
22        restartPolicy: OnFailure
```



Labels

- Labels
 - Key value pairs attached to objects
 - To specify identifying attributes of objects
 - To organize and select subset of objects
 - To query objects efficiently (cli as well gui monitoring tools)
 - Attached at creation time and can be added / modified at any time
 - Label key must be unique per object
- Example labels:
 - "release" : "stable", "release" : "canary"
 - "environment" : "dev", "environment" : "qa", "environment" : "production"
 - "tier" : "frontend", "tier" : "backend", "tier" : "cache"
 - "partition" : "customerA", "partition" : "customerB"
 - "track" : "daily", "track" : "weekly"



Labels

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-labels
5   labels:
6     environment: production
7     app: nginx
8 spec:
9   containers:
10  - name: nginx
11    image: nginx
12    ports:
13      - containerPort: 80
~
```



Selectors

- Equality Based

- environment = production
- tier != frontend

- Set Based

- environment in (production, qa)
- tier notin (frontend, backend)
- partition
- !partition



Selectors - Examples

- `kubectl get pods -l environment=production,tier=frontend`
- `kubectl get pods -l 'environment in (production),tier in (frontend)'`
- `kubectl get pods -l 'environment in (production, qa)'`
- `kubectl get pods -l 'environment,environment notin (frontend)'`

Jobs, Deployments, ReplicaSet, Daemonset

```
selector:  
  matchLabels:  
    component: redis  
  matchExpressions:  
    - {key: tier, operator: In, values: [cache]}  
    - {key: environment, operator: NotIn, values: [dev]}
```



Selectors - Examples

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: selector-pod-node
5 spec:
6   containers:
7     - name: cuda-test
8       image: "k8s.gcr.io/cuda-vector-add:v0.1"
9       resources:
10         limits:
11           nvidia.com/gpu: 1
12   nodeSelector:
13     accelerator: nvidia-tesla-p100
```



Annotations

- To attach non-identifying arbitrary metadata to objects
- Usage
 - Pointers for debugging purposes
 - Build, release, image hashes etc
 - Author info, contact details
 - Metadata to help tools for deployment, management, introspection

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: annotations-pod
5   annotations:
6     imageregistry: "https://hub.docker.com/"
7 spec:
8   containers:
9     - name: nginx
10    image: nginx
11    ports:
12      - containerPort: 80
```



Annotations

```
training > app > weather > ! 08c-metrics-weather-services.yaml
16  # microservices components 2 - api services
17  apiVersion: apps/v1
18  kind: Deployment
19  metadata:
20    labels:
21      app: weather-services
22    name: weather-services
23    namespace: weather
24  spec:
25    replicas: 1
26    selector:
27      matchLabels:
28        app: weather-services
29    template:
30      metadata:
31        labels:
32          app: weather-services
33        annotations:
34          prometheus.io/scrape: "true"
35          prometheus.io/port: "8888"
36          prometheus.io/path: /actuator/prometheus
37    spec:
```



Scalability

- Deployments
 - Rolling Updates & Rollbacks
 - Auto scaling pods
- Demo
- Practicals



Exercises - Reference

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Deployments

Use Cases

- To rollout a set of PODs
- To declare a new set of PODs
- To rollback to an earlier version of deployment
- To scale up deployment to facilitate more load
- To pause the deployment / rollout
- To autoscale deployment when cpu usage threshold reached



Deployment - Example

```
dockerk8s > kubernetes > lab > scalability > ! scalability.yaml > ...
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    labels:
5      app: google
6      name: google
7  spec:
8    replicas: 1
9    selector:
10   matchLabels:
11     app: google
12   template:
13     metadata:
14       labels:
15         app: google
16     spec:
17       imagePullSecrets:
18         - name: regcred
19       containers:
20         - image: gcr.io/google-samples/hello-app:1.0
21           name: google
22           ports:
23             - containerPort: 8080
24           resources:
25             limits:
26               cpu: "10m"
27               memory: "50Mi"
```



Deployment - Commands

- `kubectl create deployment nginx --image=nginx:1.15 --replicas=5`
- `kubectl get deployment/nginx`
- `kubectl describe deployment/nginx`
- `kubectl rollout history deployment/nginx`
- `kubectl set image deployment/nginx nginx=nginx:1.16`
- `kubectl rollout history deployment/nginx`
- `kubectl rollout undo deployment/nginx`
- `kubectl rollout undo deployment/nginx --to-revision=2`
- `kubectl scale --replicas=10 deployment/nginx`
- `kubectl rollout pause deployment/nginx`
- `kubectl rollout status deployment/nginx`
- `kubectl rollout resume deployment/nginx`
- `kubectl autoscale deployment/nginx --min=2 --max=10`



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Networking

- CNI
- CNI Flow
- Routing:
 - Within Host, Across Hosts
- Calico - Architecture
- CoreDNS
- Benchmarks / Comparison



Networking

```
Fetched 73.5 MB in 3s (21.3 MB/s)
Selecting previously unselected package conntrack.
(Reading database ... 64018 files and directories currently installed.)
Preparing to unpack .../0-conntrack_1%3a1.4.5-2_amd64.deb ...
Unpacking conntrack (1:1.4.5-2) ...
Selecting previously unselected package cri-tools.
Preparing to unpack .../1-cri-tools_1.19.0-00_amd64.deb ...
Unpacking cri-tools (1.19.0-00) ...
Selecting previously unselected package ebttables.
Preparing to unpack .../2-ebtables_2.0.11-3build1_amd64.deb ...
Unpacking ebttables (2.0.11-3build1) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../3-kubernetes-cni_0.8.7-00_amd64.deb ...
Unpacking kubernetes-cni (0.8.7-00) ...
Selecting previously unselected package socat.
Preparing to unpack .../4-socat_1.7.3.3-2_amd64.deb ...
Unpacking socat (1.7.3.3-2) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../5-kubelet_1.22.3-00_amd64.deb ...
Unpacking kubelet (1.22.3-00) ...
Selecting previously unselected package kubectl.
Preparing to unpack .../6-kubectl_1.22.3-00_amd64.deb ...
Unpacking kubectl (1.22.3-00) ...
Selecting previously unselected package kubeadm.
Preparing to unpack .../7-kubeadm_1.22.3-00_amd64.deb ...
Unpacking kubeadm (1.22.3-00) ...
Setting up conntrack (1:1.4.5-2) ...
Setting up kubectl (1.22.3-00) ...
Setting up ebttables (2.0.11-3build1) ...
Setting up socat (1.7.3.3-2) ...
Setting up cri-tools (1.19.0-00) ...
Setting up kubernetes-cni (0.8.7-00) ...
Setting up kubelet (1.22.3-00) ...
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /lib/systemd/system/kubelet.service.
Setting up kubeadm (1.22.3-00) ...
Processing triggers for man-db (2.9.1-1) ...
```



Networking

```
unigps@ip-172-31-107-158:~$ k get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd69978-56f5h	0/1	Pending	0	92s
kube-system	coredns-78fcd69978-b7fb4	0/1	Pending	0	92s
kube-system	etcd-ip-172-31-107-158	1/1	Running	0	100s
kube-system	kube-apiserver-ip-172-31-107-158	1/1	Running	0	97s
kube-system	kube-controller-manager-ip-172-31-107-158	1/1	Running	0	97s
kube-system	kube-proxy-2d8nx	1/1	Running	0	93s
kube-system	kube-scheduler-ip-172-31-107-158	1/1	Running	0	97s

```
unigps@ip-172-31-107-158:~$ █
```



Networking

```
root@ip-172-31-107-158:/etc/kubernetes# service kubelet status
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
             └─10-kubeadm.conf
     Active: active (running) since Thu 2021-11-11 05:23:22 UTC; 16min ago
       Docs: https://kubernetes.io/docs/home/
 Main PID: 4945 (kubelet)
    Tasks: 14 (limit: 2327)
   Memory: 40.1M
      CGroup: /system.slice/kubelet.service
              └─4945 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml --network-plugin=cni -->

Nov 11 05:39:22 ip-172-31-107-158 kubelet[4945]: I1111 05:39:22.678325 4945 cni.go:239] "Unable to update cni config" err="no networks found in /etc/cni/net.d"
Nov 11 05:39:25 ip-172-31-107-158 kubelet[4945]: E1111 05:39:25.556760 4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:27 ip-172-31-107-158 kubelet[4945]: I1111 05:39:27.678493 4945 cni.go:239] "Unable to update cni config" err="no networks found in /etc/cni/net.d"
4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:30 ip-172-31-107-158 kubelet[4945]: E1111 05:39:30.567582 4945 cni.go:239] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:32 ip-172-31-107-158 kubelet[4945]: I1111 05:39:32.679221 4945 cni.go:239] "Unable to update cni config" err="no networks found in /etc/cni/net.d"
4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:35 ip-172-31-107-158 kubelet[4945]: E1111 05:39:35.578361 4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:37 ip-172-31-107-158 kubelet[4945]: I1111 05:39:37.679897 4945 cni.go:239] "Unable to update cni config" err="no networks found in /etc/cni/net.d"
4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:40 ip-172-31-107-158 kubelet[4945]: E1111 05:39:40.588890 4945 cni.go:239] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:42 ip-172-31-107-158 kubelet[4945]: I1111 05:39:42.680298 4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
Nov 11 05:39:45 ip-172-31-107-158 kubelet[4945]: E1111 05:39:45.599653 4945 cni.go:239] "Unable to update cni config" err="no networks found in /etc/cni/net.d"
4945 kubelet.go:2337] "Container runtime network not ready" networkReady=false reason:NetworkPluginNotReady message:>
```



Networking

```
root@ip-172-31-107-158:/etc/kubernetes# docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
a7099bc9d43e        6120bd723dce   "/usr/local/bin/kube..."   19 minutes ago    Up 19 minutes
49a813743c95        k8s.gcr.io/pause:3.5  "/pause"
3363f0d3ca15        53224b502ea4   "kube-apiserver --ad..."  19 minutes ago    Up 19 minutes
0
5357e5614358        004811815584   "etcd --advertise-cl..."
7930bffe223d        0aa9c7e31d30   "kube-scheduler --au..."
0
ceee5e692ebd        05c905cef780   "kube-controller-man..."
3271cf6166acc0755_0
26bebcdcc57a        k8s.gcr.io/pause:3.5  "/pause"
934f2af6644b        k8s.gcr.io/pause:3.5  "/pause"
60bc7ca7c123        k8s.gcr.io/pause:3.5  "/pause"
aed4d7056f21        k8s.gcr.io/pause:3.5  "/pause"
root@ip-172-31-107-158:/etc/kubernetes#
```



CNI - Interface Methods

- Add
 - Add a container to the network
- DEL
 - Delete a container from the network
- CHECK
 - Return an error if any issue
- VERSION
 - Report version info about the plugin

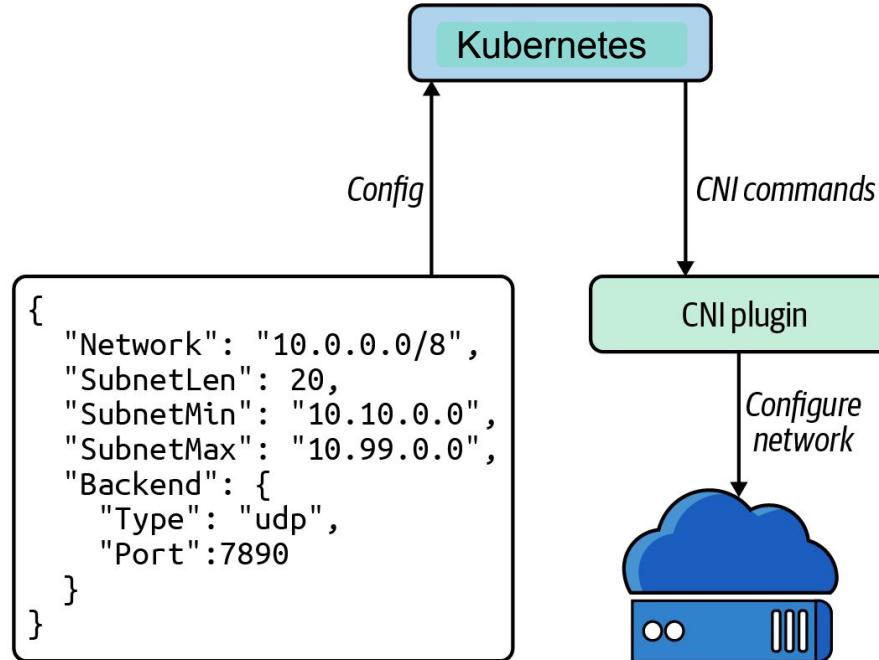
Ref: <https://github.com/containernetworking/cni/blob/master/SPEC.md>



CNI - Flow

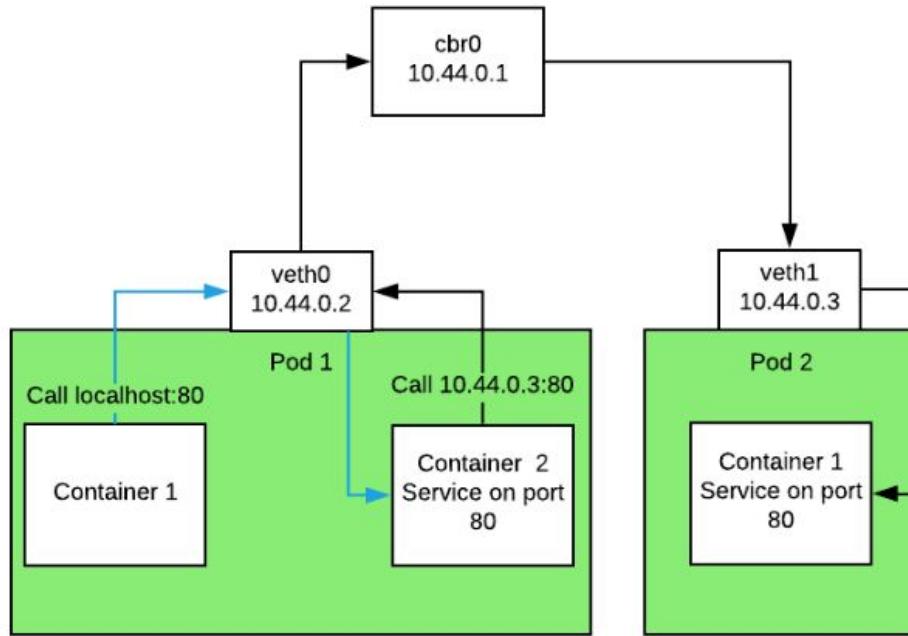
CNI Requirements

- IPAM - IP Address Management
- Connectivity in Container Network
- Route Advertisement



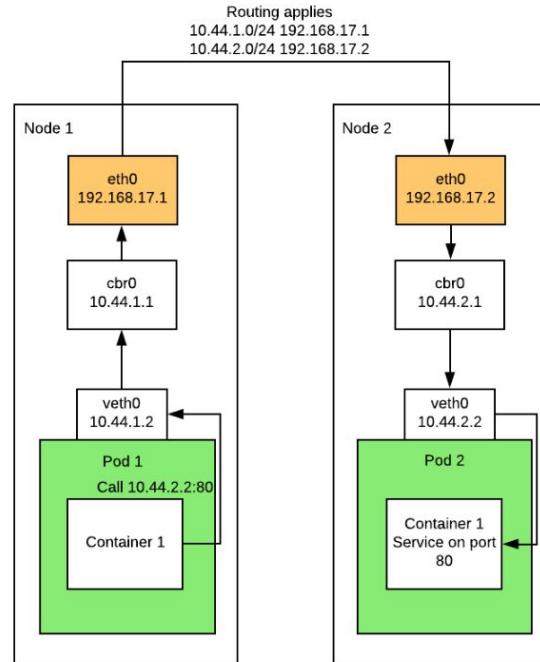


Network Routing - Same Host





Network Routing - Across Nodes

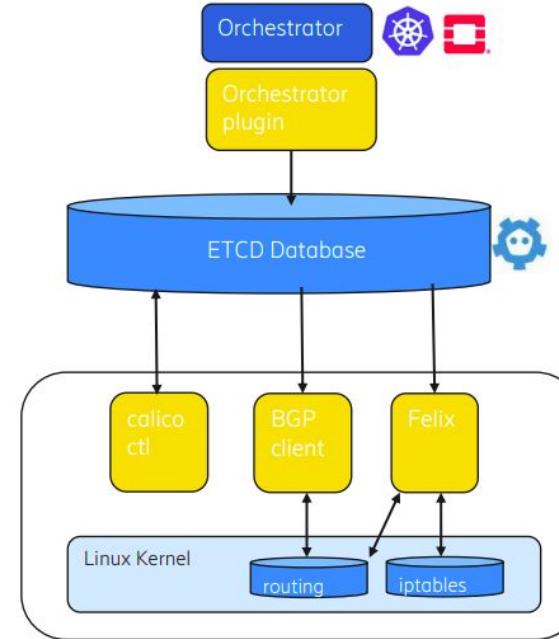




Calico Architecture

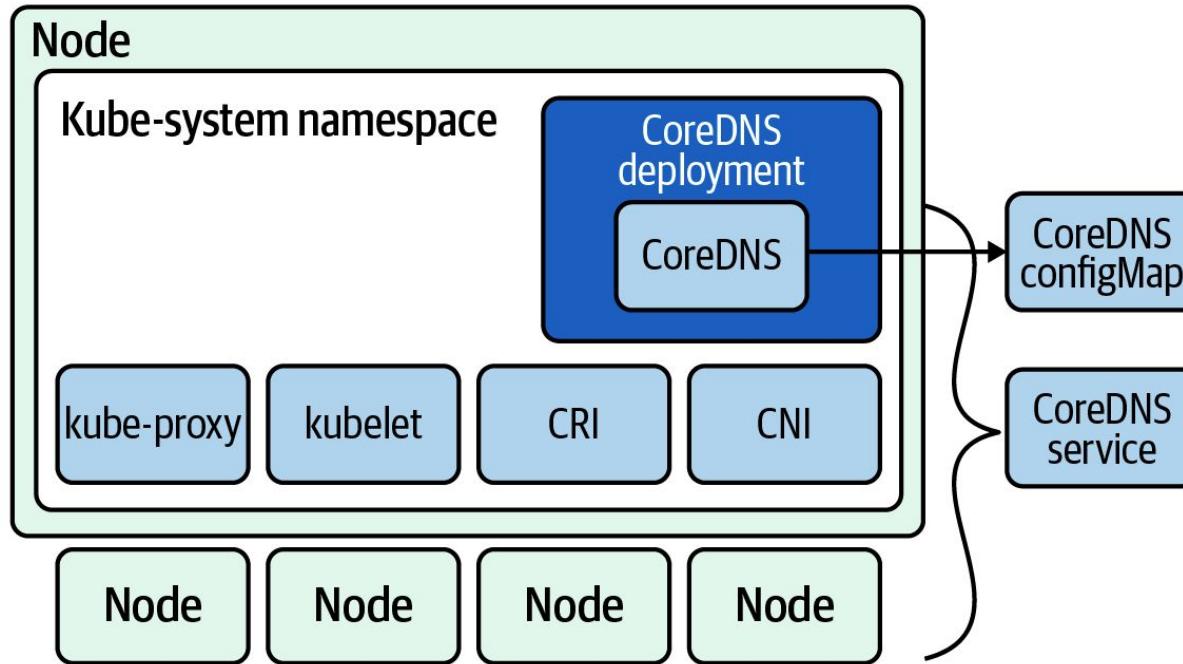
Features

- Layer 3 Routing
- BGP for Routes distribution
- Support Network Policy
- Uses iptables & kernel routing





Networking - CoreDNS



Ref: Oreilly docs



CNI Plugins - Benchmarking

CNI Benchmark August 2020 infraBuilder	Config	Performances (bandwidth)				Resources consumption (cpu/ram)				Security features					
		MTU	Pod to Pod		Pod to Service		Idle	Pod to Pod		Pod to Service		Network Policies	Encryption		
	setting	setting	TCP	UDP	TCP	UDP	none	TCP	UDP	TCP	UDP	in	out	activation	Performance
Antrea	auto	Very fast	Very fast	Very fast	Slow	Low	Low	Low	Low	Low	Low	yes	yes	at deploy time	Slow
Calico	manual	Very fast	Very fast	Very fast	Fast	Low	Very low	Very low	Very low	Very low	Very low	yes	yes	anytime	Very fast
Canal	manual	Very fast	Very fast	Very fast	Very fast	Low	Very low	Very low	Very low	Very low	Very low	yes	yes	no	n/a
Cilium	auto	Fast	Very fast	Very fast	Very fast	High	High	High	High	High	High	yes	yes	at deploy time	Slow
Flannel	auto	Very fast	Very fast	Very fast	Very fast	Very low	Very low	Very low	Very low	Very low	Very low	no	no	no	n/a
Kube-OVN	auto	Fast	Very slow	Fast	Very slow	High	High	High	High	High	High	yes	yes	no	n/a
Kube-router	none	Slow	Very slow	Slow	Very slow	Low	Very low	Low	Very low	Low	Low	yes	yes	no	n/a
Weave Net	manual	Very fast	Very fast	Very fast	Fast	Very low	Low	Low	Low	Low	Low	yes	yes	at deploy time	Slow



Volumes | Persistence

- Volumes
- Persistence Volume (pv)
- Persistence Volume Claim (pvc)
- Statefulset (sts)
- Daemonset (ds)



Overview - Volumes | Persistent

- Ephemeral Volumes
 - Tightly coupled with POD lifetime
 - Deleted when POD is removed
 - Example: emptydir
- Persistent Volumes
 - Survives POD reboots
 - Meant for long term and independent of POD / Node lifecycle
 - Examples: hostpath, NFS, Cloud storage (EBS etc)
- The access modes are:
 - ReadWriteOnce – the volume can be mounted as read-write by a single node
 - ReadOnlyMany – the volume can be mounted read-only by many nodes
 - ReadWriteMany -- the volume can be mounted as read-write by many nodes



Volume - emptyDir

training > kubernetes > lab > persistence > ! emptyDir-volume.yaml

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: volume-emptydir
5  spec:
6    containers:
7      - image: nginx
8        name: test-container
9        volumeMounts:
10          - mountPath: /cache
11            name: cache-volume
12    volumes:
13      - name: cache-volume
14        emptyDir: {}
```



PV - hostpath

```
sljpmk8s > kubernetes > lab > persistence > ! hostpath-volume.yaml >
          io.k8s.api.core.v1.Pod (v1@pod.json)
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: hostpath-volume
5  spec:
6    containers:
7      - image: nginx
8        name: test-container
9        volumeMounts:
10       - mountPath: /data-mounted-as
11         name: hostpath-volume
12         resources:
13           limits:
14             cpu: "20m"
15             memory: "50Mi"
16         volumes:
17           - name: hostpath-volume
18             hostPath:
19               # directory location on host
20               path: /data
21               # this field is optional
22               type: DirectoryOrCreate
```



PV - Cloud Storage (AWS EBS)

```
sljpmk8s > kubernetes > lab > persistence > ! 03-aws-ebs-volume.yaml >
          io.k8s.api.core.v1.Pod (v1@pod.json)
1   apiVersion: v1
2   kind: Pod
3   metadata:
4     name: volume-ebs
5   spec:
6     nodeName: ip-172-31-58-74.ec2.internal
7     containers:
8       - image: nginx
9         name: test-container
10        volumeMounts:
11          - mountPath: /test-ebs
12            name: ebs-volume
13        resources:
14          limits:
15            cpu: "20m"
16            memory: "50Mi"
17        volumes:
18          - name: ebs-volume
19            # This AWS EBS volume must already exist.
20            awsElasticBlockStore:
21              volumeID: vol-081b401a592b794a2
22              fsType: ext4
```

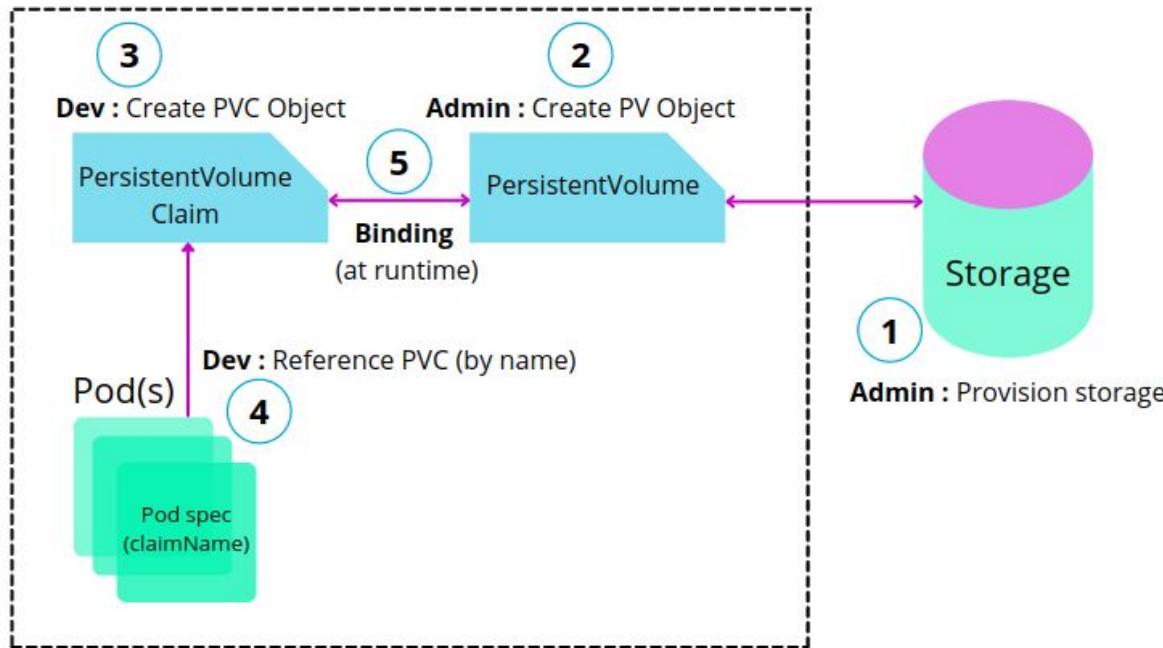


PV - Config Volume

```
sljpmk8s > kubernetes > lab > persistence > ! 04-config-volume.yaml >
          io.k8s.api.core.v1.Pod (v1@pod.json) | io.k8s.api.core.v1.ConfigMap (v10
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: config-volume
5  spec:
6    containers:
7      - name: test
8        image: nginx
9        volumeMounts:
10       - name: config-vol
11         mountPath: /etc/config
12    resources:
13      limits:
14        cpu: "20m"
15        memory: "50Mi"
16    volumes:
17      - name: config-vol
18        configMap:
19          name: config-volume
20          items:
21            - key: log_level
22              path: log.properties
```

```
24  apiVersion: v1
25  kind: ConfigMap
26  metadata:
27    name: config-volume
28  data:
29    log_level: |
30      debug=true
```

Persistent Volume - static





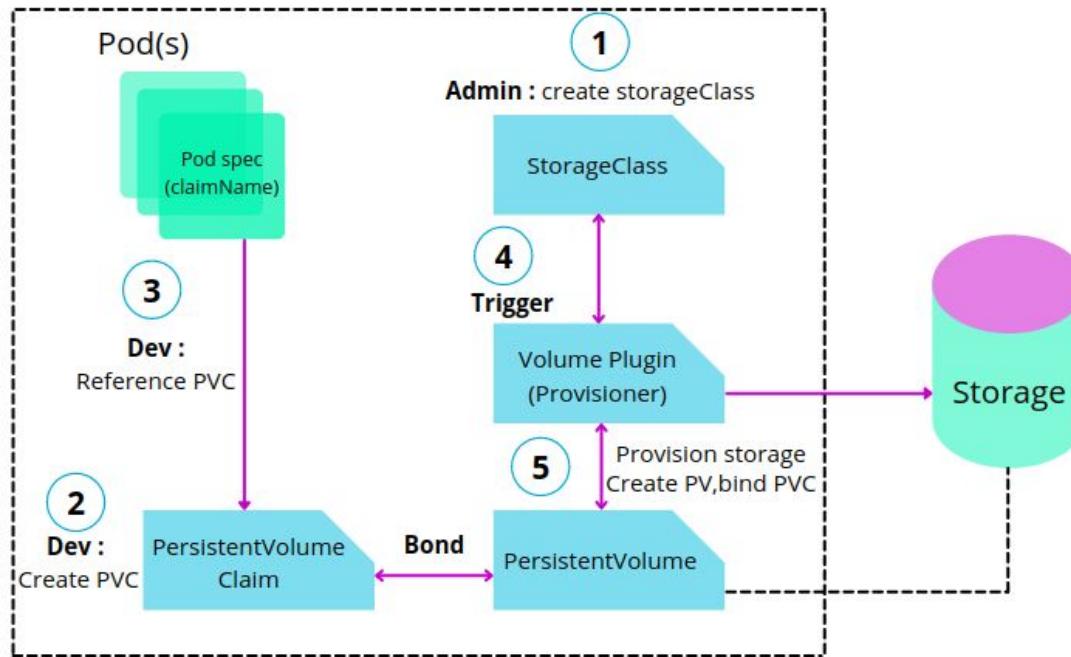
PV - Static

```
1 apiVersion: v1
2 kind: PersistentVolume
3 metadata:
4   name: pv-hostpath
5   annotations:
6     pv.beta.kubernetes.io/gid: "1234"
7   labels:
8     type: local
9 spec:
10  storageClassName: manual
11  capacity:
12    storage: 10Gi
13  accessModes:
14    - ReadWriteOnce
15  hostPath:
16    path: "/mnt/data"
```

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-test
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10  requests:
11    storage: 3Gi
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-pv-pvc
5 spec:
6   volumes:
7     - name: pv-storage
8       persistentVolumeClaim:
9         claimName: pvc-test
10 containers:
11   - name: task-pv-container
12     image: nginx
13     ports:
14       - containerPort: 80
15         name: "http-server"
16     volumeMounts:
17       - mountPath: "/usr/share/nginx/html"
18         name: pv-storage
```

Persistent Volume - Dynamic





Example

```
1 kind: StorageClass
2 apiVersion: storage.k8s.io/v1
3 metadata:
4   name: storageclass-generic
5 provisioner: kubernetes.io/aws-ebs
6 parameters:
7   type: gp2
8   zones: us-east-1a,us-east-1b,us-east-1c
9   iopsPerGB: "10"
10  fsType: ext4
```

```
1 apiVersion: v1
2 kind: PersistentVolumeClaim
3 metadata:
4   name: pvc-dynamic
5   labels:
6     app: nginx
7 spec:
8   storageClassName: storageclass-generic
9   accessModes:
10    - ReadWriteOnce
11   resources:
12     requests:
13       storage: 1Gi
```

```
1 kind: Pod
2 apiVersion: v1
3 metadata:
4   name: pod-volume-dynamic
5   labels:
6     app: nginx
7 spec:
8   containers:
9     - name: nginx
10    image: nginx
11   volumeMounts:
12    - mountPath: "/var/www/html"
13      name: external
14 volumes:
15  - name: external
16    persistentVolumeClaim:
17      claimName: pvc-dynamic
```



Persistent Volumes

- GCEPersistentDisk
- AWSElasticBlockStore
- AzureFile
- AzureDisk
- CSI
- FC (Fibre Channel)
- FlexVolume
- Flocker
- NFS
- iSCSI
- RBD (Ceph Block Device)
- CephFS
- Cinder (OpenStack block storage)
- Glusterfs
- VsphereVolume
- Quobyte Volumes
- HostPath (Single node testing only – local storage is not supported in any way and WILL NOT WORK in a multi-node cluster)
- Portworx Volumes
- ScaleIO Volumes
- StorageOS



StatefulSet

Use Cases

- Stable, unique network identifiers
- Stable, persistent storage
- Ordered, graceful deployment and scaling
- Ordered, automated rolling updates

Limitations

- No automatic deletion of referenced volumes
- No PODs deletion guarantee when StatefulSet is deleted
- Rolling Updates not consistent always



StatefulSet - Example

```
1 apiVersion: apps/v1
2 kind: StatefulSet
3 metadata:
4   name: sts-web
5 spec:
6   serviceName: "nginx"
7   replicas: 2
8   selector:
9     matchLabels:
10    app: nginx
11 template:
12   metadata:
13     labels:
14       app: nginx
15   spec:
16     containers:
17       - name: nginx
18         image: k8s.gcr.io/nginx-slim:0.8
19         ports:
20           - containerPort: 80
21             name: web
22         volumeMounts:
23           - name: www
24             mountPath: /usr/share/nginx/html
25   volumeClaimTemplates:
26     - metadata:
27       name: www
28     spec:
29       accessModes: [ "ReadWriteOnce" ]
30       resources:
31         requests:
32           storage: 1Gi
```

- Scale Up
- Scale down
- `kubectl set image sts/sts-web nginx=nginx:1.18`
- Staged Update
 - `kubectl patch statefulset sts-web -p '{"spec":{"updateStrategy":{"type":"RollingUpdate","rollingUpdate":{"partition":3}}}}'`
 - `kubectl patch statefulset sts-web --type='json' -p='[{"op": "replace", "path": "/spec/template/spec/containers/0/image", "value":"nginx:1.17"}]'`
 - `kubectl get pod sts-web-1 --template '{{range $i, $c := .spec.containers}}{{{$c.image}}}{{end}}'`



DaemonSet

Purpose

- To run a copy of a POD on all / some node(s)

Use Cases

- Storage cluster daemon (gluster, ceph)
- Log Collectors (fluentd, logstash)
- Node Monitoring daemons (Prometheus, Dynatrace, collectd)



DaemonSet - Example

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-elasticsearch
  namespace: kube-system
  labels:
    k8s-app: fluentd-logging
spec:
  selector:
    matchLabels:
      name: fluentd-elasticsearch
  template:
    metadata:
      labels:
        name: fluentd-elasticsearch
    spec:
      tolerations:
      - key: node-role.kubernetes.io/master
        effect: NoSchedule
      containers:
      - name: fluentd-elasticsearch
        image: quay.io/fluentd_elasticsearch/fluentd:v2.5.2
        resources:
          limits:
            memory: 200Mi
          requests:
            cpu: 100m
            memory: 200Mi
        volumeMounts:
        - name: varlog
          mountPath: /var/log
        - name: varlibdockercontainers
          mountPath: /var/lib/docker/containers
          readOnly: true
      terminationGracePeriodSeconds: 30
  volumes:
  - name: varlog
    hostPath:
      path: /var/log
  - name: varlibdockercontainers
    hostPath:
      path: /var/lib/docker/containers
```



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Configuration

- Config Maps
- Secrets
- Security Contexts
- Accounts
- Demo
- Practicals



ConfigMap

- To store non-confidential key-value pairs
- Can be consumed as env variables, command line args or config files in volume
- To decouple env specific config from images for portability
- Max data 1MB

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: cm-game-demo
5 data:
6   # property-like keys; each key maps to a simple value
7   player_initial_lives: "3"
8   ui_properties_file_name: "user-interface.properties"
9   # file-like keys
10  game.properties: |
11    enemy.types=aliens,monsters
12    player.maximum-lives=5
13  user-interface.properties: |
14    color.good=purple
15    color.bad=yellow
16    allow.textmode=true
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-configmap
5 spec:
6   containers:
7     - name: demo
8       image: alpine
9       command: ["sleep", "3600"]
10      env:
11        # Define the environment variable
12        - name: PLAYER_INITIAL_LIVES # Notice that the case is different here
13          valueFrom:
14            configMapKeyRef:
15              name: cm-game-demo
16              key: player_initial_lives # The ConfigMap this value comes from.
17            - name: UI_PROPERTIES_FILE_NAME
18              valueFrom:
19                configMapKeyRef:
20                  name: CM-game-demo
21                  key: ui_properties_file_name
22            volumeMounts:
23              - name: config
24                mountPath: "/config"
25                readOnly: true
26    volumes:
27      # You set volumes at the Pod level, then mount them into containers inside that Pod
28      - name: config
29        configMap:
30          # Provide the name of the ConfigMap you want to mount.
31          name: cm-game-demo
32          # An array of keys from the ConfigMap to create as files
33          items:
34            - key: "game.properties"
35            path: "game.properties"
36            - key: "user-interface.properties"
37            path: "user-interface.properties"
```



Secret

- To manage sensitive info like password, oauthkeys, docker login, ssh keys, tls etc
- Examples

```
kubectl create secret docker-registry secret-tiger-docker \
    --docker-username=tiger \
    --docker-password=pass113 \
    --docker-email=tiger@acme.com
```

```
kubectl create secret tls my-tls-secret \
    --cert=path/to/cert/file \
    --key=path/to/key/file
```

Secret

```
1 apiVersion: v1
2 data:
3   username: YWRtaW4=
4   password: MWYyZDFlMmU2N2Rm
5 kind: Secret
6 metadata:
7   name: pod-secret
8   namespace: default
9   resourceVersion: "164619"
10  uid: cfee02d6-c137-11e5-8d73-42010af00002
11 type: Opaque
```

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: pod-secret
5 spec:
6   containers:
7     - name: mypod
8       image: redis
9       env:
10      - name: SECRET_USERNAME
11        valueFrom:
12          secretKeyRef:
13            name: mysecret
14            key: username
15      - name: SECRET_PASSWORD
16        valueFrom:
17          secretKeyRef:
18            name: mysecret
19            key: password
20   volumeMounts:
21     - name: foo
22       mountPath: "/etc/foo"
23       readOnly: true
24   volumes:
25     - name: foo
26       secret:
27         secretName: pod-secret
28         items:
29           - key: username
30             path: my-group/my-username
=
```



Security Context

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: security-context-demo
5 spec:
6   securityContext:
7     runAsUser: 1000
8     runAsGroup: 3000
9     fsGroup: 2000
10  volumes:
11    - name: sec-ctx-vol
12      emptyDir: {}
13  containers:
14    - name: sec-ctx-demo
15      image: busybox
16      command: [ "sh", "-c", "sleep 1h" ]
17      volumeMounts:
18        - name: sec-ctx-vol
19          mountPath: /data/demo
20      securityContext:
21        allowPrivilegeEscalation: false
22        capabilities:
23          add: ["NET_ADMIN", "SYS_TIME"]
24
```



User Accounts & Service Accounts

- User Accounts
 - User accounts are for humans.
 - User accounts are intended to be global. Names must be unique across all namespaces of a cluster.
- Service Accounts
 - Service accounts are for processes, which run in pods.
 - Service accounts are namespaced.
 - Service account creation is intended to be more lightweight



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Scheduling

- Node Name
- Node Selector
- Affinity (Pod, Node)
- Taints & Tolerations
- Demo
- Practicals



Overview

Scheduling - Matching PODs to Nodes so that kubelet can run them

Filtering

Scoring



Node Name

```
training > kubernetes > lab > scheduling > ! schedule-nodename.yaml
 1  apiVersion: apps/v1
 2  kind: Deployment
 3  metadata:
 4    name: schedule-node
 5  spec:
 6    replicas: 1
 7    selector:
 8      matchLabels:
 9        component: schedule-node
10    template:
11      metadata:
12        labels:
13          component: schedule-node
14    spec:
15      nodeName: ip-172-31-102-18.ec2.internal
16      containers:
17        - name: test-app
18          image: brainupgrade/test-app:all-tiers-in-one
19          imagePullPolicy: IfNotPresent
20          ports:
21            - containerPort: 8080
22          resources:
23            requests:
24              cpu: "100m"
25              memory: "250Mi"
```



Node Selector

```
training > kubernetes > lab > scheduling > ! schedule-nodeselector.yaml
 1  apiVersion: apps/v1
 2  kind: Deployment
 3  metadata:
 4    name: schedule-nodeselector
 5  spec:
 6    replicas: 1
 7    selector:
 8      matchLabels:
 9        component: schedule-nodeselector
10    template:
11      metadata:
12        labels:
13          component: schedule-nodeselector
14    spec:
15      nodeSelector:
16        node-role.kubernetes.io/spot-worker: "true"
17      containers:
18        - name: test-app
19          image: nginx
20          imagePullPolicy: IfNotPresent
21          ports:
22            - containerPort: 80
```



POD - Affinity & Anti Affinity (example 1)

```
training > kubernetes > lab > scheduling > ! pod-affinity.yaml
14     spec:
15         affinity:
16             podAffinity:
17                 requiredDuringSchedulingIgnoredDuringExecution:
18                     - labelSelector:
19                         matchExpressions:
20                             - key: tier
21                                 operator: In # NotIn, Exists, NotExists
22                                     values:
23                                         - cache
24                                         topologyKey: topology.kubernetes.io/zone
25             podAntiAffinity:
26                 preferredDuringSchedulingIgnoredDuringExecution:
27                     - weight: 100
28                     podAffinityTerm:
29                         labelSelector:
30                             matchExpressions:
31                                 - key: tier
32                                     operator: In
33                                         values:
34                                         - messaging
35                                         topologyKey: kubernetes.io/hostname
36             containers:
37                 - name: pod-affinity-middle-tier
38                     image: nginx
39                     imagePullPolicy: IfNotPresent
40                     ports:
41                         - containerPort: 80
```



POD - Affinity & Anti Affinity (example 2)

```
training > kubernetes > lab > scheduling > pod-affinity-web-cache > ! app-cache-store.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: app-cache-store
5  spec:
6    selector:
7      matchLabels:
8        app: cache-store
9    replicas: 2
10   template:
11     metadata:
12       labels:
13         app: cache-store
14   spec:
15     affinity:
16       podAntiAffinity:
17         requiredDuringSchedulingIgnoredDuringExecution:
18           - labelSelector:
19             matchExpressions:
20               - key: app
21                 operator: In
22                   values:
23                     - cache-store
24           topologyKey: "kubernetes.io/hostname"
25   containers:
26     - name: redis-server
27       image: redis:3.2-alpine
```

```
training > kubernetes > lab > scheduling > pod-affinity-web-cache > ! app-web-server.yaml
9   replicas: 2
10  template:
11    metadata:
12      labels:
13        app: web-server
14    spec:
15      affinity:
16        podAntiAffinity:
17          requiredDuringSchedulingIgnoredDuringExecution:
18            - labelSelector:
19              matchExpressions:
20                - key: app
21                  operator: In
22                    values:
23                      - web-server
24          topologyKey: "kubernetes.io/hostname"
25      podAffinity:
26        requiredDuringSchedulingIgnoredDuringExecution:
27          - labelSelector:
28            matchExpressions:
29              - key: app
30                operator: In
31                  values:
32                    - cache-store
33          topologyKey: "kubernetes.io/hostname"
34      containers:
35        - name: web-app
36          image: nginx:1.16-alpine
```



Node - Affinity

```
training > kubernetes > lab > scheduling > ! node-affinity.yaml
14     spec:
15         affinity:
16             nodeAffinity:
17                 requiredDuringSchedulingIgnoredDuringExecution:
18                     nodeSelectorTerms:
19                         - matchExpressions:
20                             - key: topology.kubernetes.io/zone
21                                 operator: In #In, NotIn, Exists, DoesNotExist, Gt, Lt
22                                 values:
23                                     - us-east-1b
24                                     - us-east-1c
25                 preferredDuringSchedulingIgnoredDuringExecution:
26                     - weight: 1
27                         preference:
28                             matchExpressions:
29                                 - key: kops.k8s.io/instancegroup
30                                     operator: In #NotIn, DoesNotExist for AntiAffinity
31                                     values:
32                                         - spotnodes
33             containers:
34                 - name: nginx
35                   image: nginx
36                   imagePullPolicy: IfNotPresent
37                   ports:
38                       - containerPort: 80
```



Taints & Tolerations

- Objective

To steer pods away from nodes or evict pods that should not be running on particular nodes

- Use Cases

- Dedicated Nodes
- Nodes with special hardware

- Notes

- Taints - Node (to reject pods that don't tolerate taints)
- Tolerations - Pod (to allow pods to schedule onto nodes matching taints)



Taints

```
lab@rajesh-Gazelle:~$ kubectl get nodes -o json | jq ".items[]|{name:.metadata.name, taints:.spec.taints}"  
{  
  "name": "ip-172-31-102-18.ec2.internal",  
  "taints": null  
}  
{  
  "name": "ip-172-31-34-27.ec2.internal",  
  "taints": [  
    {  
      "effect": "NoSchedule",  
      "key": "node-role.kubernetes.io/master"  
    }  
  ]  
}  
{  
  "name": "ip-172-31-87-231.ec2.internal",  
  "taints": null  
}
```

Effect types: NoSchedule or PreferNoSchedule or NoExecute



Taints

- Examples (Taints)

- To apply

```
kubectl taint nodes qanode release=qa:NoSchedule
```

```
kubectl taint nodes devnode release=dev:PreferNoSchedule
```

```
kubectl taint nodes prodnodet release=prod:NoExecute
```

- To remove

```
kubectl taint nodes node-name key=value:Effect-
```



Tolerations

```
training > kubernetes > lab > scheduling > ! tolerations-master.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: tolerations-master
5  spec:
6    containers:
7      - image: nginx
8        name: nginx
9    dnsPolicy: ClusterFirst
10   restartPolicy: Always
11   tolerations:
12     - key: "node-role.kubernetes.io/master"
13       operator: "Exists"
14       effect: "NoSchedule"
```



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Services

- Cluster IP
- Node Port
- Load Balancer
- External Name
- Demo
- Practicals



Service

- An abstract way to expose an application running on pod as network service.
- Frontends and backends of application can connect without worrying about POD IPs
- Uses session affinity while connecting to backend PODs

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   labels:
5     app: nginx
6   name: nginx
7 spec:
8   ports:
9     - port: 80
10    protocol: TCP
11    targetPort: 80
12 selector:
13   app: nginx
```



Service Types

- Cluster IP
 - Service exposed on cluster internal IP
 - Reachable only within cluster
- Node Port
 - Exposed on each Node IP at static port (30000-32767)
- Load Balancer
 - Exposed through external cloud load balancer
- ExternalName
 - Maps the service to the external name like api.example.com



Service - External Name

```
training > kubernetes > lab > services > ! externalName.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: external-name
5  spec:
6    type: ExternalName
7    externalName: google.com
```



Service - External IP

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 9376
  externalIPs:
    - 80.11.12.10
```

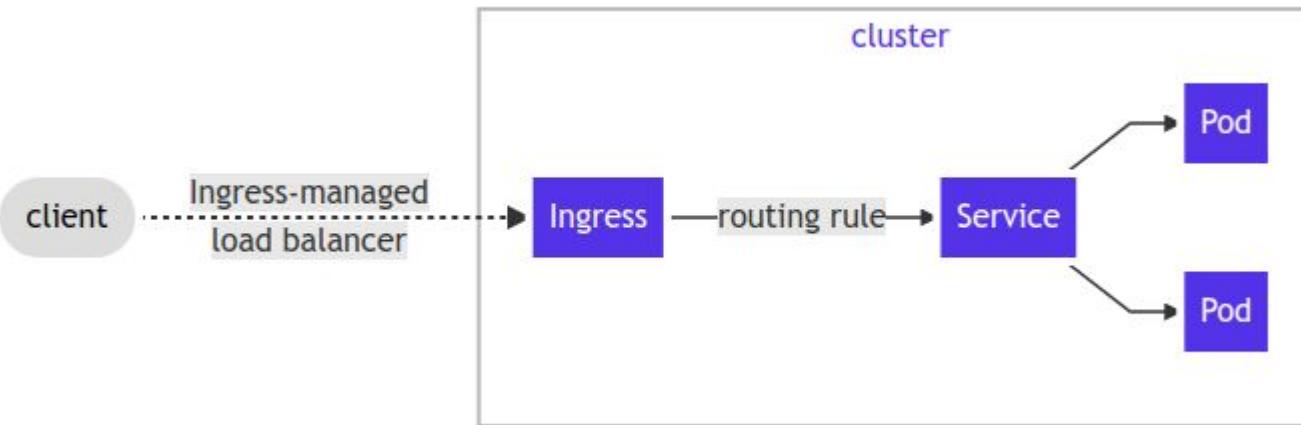
Use Cases:

- External DB Cluster in production
- To point a service in another namespace / cluster



Ingress

- Provides load balancing, SSL Termination and Name based virtual hosting
- Provides externally reachable URLs to Services
- Used for HTTP / HTTPS protocols



Source: kubernetes.io



Ingress

```
training > kubernetes > lab > ingress > ! ingress.yaml
  1  apiVersion: networking.k8s.io/v1beta1
  2  kind: Ingress
  3  metadata:
  4    name: rajeshg.brainupgrade.in
  5    annotations:
  6      cert-manager.io/cluster-issuer: letsencrypt-prod
  7      kubernetes.io/ingress.class: nginx
  8      kubernetes.io/tls-acme: "true"
  9      nginx.ingress.kubernetes.io/enable-cors: "true"
 10     nginx.ingress.kubernetes.io/proxy-body-size: 50m
 11     nginx.ingress.kubernetes.io/proxy-read-timeout: "3600"
 12     nginx.ingress.kubernetes.io/rewrite-target: /
 13     nginx.ingress.kubernetes.io/use-regex: "true"
 14   spec:
 15     tls:
 16       - hosts:
 17         - rajeshg.brainupgrade.in
 18         secretName: rajeshg.brainupgrade.in
 19       rules:
 20         - host: rajeshg.brainupgrade.in
 21           http:
 22             paths:
 23               - path: /?(.*)
 24                 backend:
 25                   serviceName: api-v1
 26                   servicePort: 8080
 27               - path: /api/?(.*)
 28                 backend:
 29                   serviceName: api-v2
 30                   servicePort: 8080
 31               - path: /logger/?(.*)
 32                 backend:
 33                   serviceName: logger
 34                   servicePort: 80
```



Network Policy

```
training > kubernetes > lab > networking > ! network-policy.yaml
 1  apiVersion: networking.k8s.io/v1
 2  kind: NetworkPolicy
 3  metadata:
 4    name: weather-network-policy
 5    namespace: weather
 6  spec:
 7    podSelector:
 8      matchLabels:
 9        app: weather-db
10    policyTypes:
11      - Ingress
12      - Egress
13    ingress:
14      - from:
15        # - ipBlock:
16        #   cidr: 172.17.0.0/16
17        #   except:
18        #     - 172.17.1.0/24
19        - namespaceSelector:
20          matchLabels:
21            app: weather
22        - podSelector:
23          matchLabels:
24            app: weather-service
25        ports:
26          - protocol: TCP
27            port: 3306
28    egress:
29      - to:
30        - ipBlock:
31          cidr: 10.0.0.0/24
32        ports:
33          - protocol: TCP
34            port: 5978
```



Examples

- Nginx deployment, Expose as service and publish (Demo)
- Spring Boot with DB (walkthrough)
- API Services rollout (walkthrough)



Exercises

Please refer the google classwork link given in the chat message

And do all the lab work as per the instructions noted in the classwork assignments



Cluster Upgrade - KubeAdm

Master Nodes

- apt update && kubeadm upgrade apply
- apt-cache madison kubeadm
- apt-mark unhold kubeadm && apt-get update && apt-get install -y kubeadm=1.18.5-00 && apt-mark hold kubeadm
- kubeadm version
- kubectl drain <cp-node-name> --ignore-daemonsets
- sudo kubeadm upgrade apply v1.18.x
- kubectl uncordon <cp-node-name>
- kubectl drain <cp-node-name> --ignore-daemonsets
- kubeadm upgrade node
- kubectl uncordon <cp-node-name>

Worker Nodes

- apt-mark unhold kubelet kubectl && apt-get update && apt-get install -y kubelet=1.18.x-00 kubectl=1.18.x-00 && apt-mark hold kubelet kubectl
- sudo systemctl daemon-reload
- sudo systemctl restart kubelet
- apt-mark unhold kubeadm && apt-get update && apt-get install -y kubeadm=1.18.x-00 && apt-mark hold kubeadm
- kubectl drain <node-to-drain> --ignore-daemonsets
- kubeadm upgrade node
- apt-mark unhold kubelet && apt-get update && apt-get install -y kubelet=1.18.x-00 && apt-mark hold kubelet
- kubectl uncordon <node-to-drain>



Helm

- What it is
 - Kubernetes package manager, collection of templates & settings also known as helm charts
- Benefits
 - Smooth deployment of sophisticated applications
 - Especially useful when app to be deployed across many clusters with custom configuration
 - Easy upgrade and rollback



Helm - Package Publish Deploy

Package & Publish

- Create git repo named helm-charts and use branch gh-pages
- Create helm chart (kubernetes application package)
 - helm create my-nginx
- Package the app
 - helm package my-nginx
- Create helm repo index from outside of git directory
 - helm repo index helm-charts --url https://<git-username>.github.io/helm-charts
- Add git changes, commit and push the changes to gh-pages
 - git add .
 - git commit -m "my helm chart" .
 - git push origin gh-pages

Deploy

- helm repo add my-charts https://<git-username>.github.io/helm-charts
- helm repo update
- helm install my-nginx my-charts/my-nginx

Upgrade & Rollback

- helm upgrade my-nginx my-nginx
- helm rollback my-nginx 1



Monitoring & Misc

- Kubernetes Dashboard
- Port Forwarding
- Introspection & Debugging
- Container Logging
- Best Practices
- Demo
- Practicals



Monitoring Dashboard (Prometheus-Grafana)





Port Forwarding

```
kubectl port-forward <pod> 7000:6379
```

```
kubectl port-forward <deployment> 7000:6379
```

```
kubectl port-forward <svc> 7000:6379
```

To access LoadBalancer service on localhost

```
minikube tunnel
```



Introspection & debugging

```
kubectl get pods  
kubectl get pod <pod-name> -o yaml  
kubectl describe <pod-name>  
kubectl describe <pod-name> -o yaml  
kubectl get events  
kubectl get events --namespace=my-namespace  (--all-namespaces)  
kubectl get nodes  
kubectl get node <node-name>  
kubectl get node <node-name> -o yaml  
kubectl describe node <node-name>  
kubectl describe node <node-name> -o yaml  
kubectl logs -f pod <pod-name>
```



Shell to running container

```
rajesh@rajesh-Gazelle:~/git/kubernetes/debugging/shell$ kubectl apply -f shell-demo.yaml

kubectl get pod shell-demo
kubectl exec -it shell-demo -- /bin/bash

root@shell-demo:/# ls /

root@shell-demo:/# echo Hello shell demo > /usr/share/nginx/html/index.html
root@shell-demo:/# apt-get update
root@shell-demo:/# apt-get install curl
root@shell-demo:/# curl localhost

kubectl exec shell-demo env

kubectl exec -it my-pod --container main-app -- /bin/bash
```



Best Practices

- Configuration - specify latest stable API version
- Keep config files in version control before pushing to cluster
- Prefer YAML over JSON
- Group related objects into one file whenever it makes more sense
- Don't specify default values unnecessarily
- Put Object descriptions as part of annotations
- Don't use naked PODs
- Avoid using hostPort for POD
- Use labels effectively
- Use image tag instead of using latest as the default
- Use kubectl run and expose to launch single container deployments & services



Tips for CKAD & CKA (CNCF)

Preparation

- Improve typing speed (Ref: <https://typing.com>)
- Improve VI skills (Ref: <https://www.youtube.com/watch?v=5r6yzFEXajQ>)
- Once concepts are clear then practice practice & practice (Get your linux os and install minikube or microk8s to play with the cluster locally)
- Question Difficulty levels: Mix of easy, intermediate and difficult (however most of the Qs are intermediate or less difficult)

Exam day

- Adequate sleep previous day
- Choose quiet place for exam
- Ensure your body is hydrated enough
- Keep yourself calm & maintain it throughout exam period

Resources

- Refer Kubernetes Documentation (<https://kubernetes.io/docs/home/>)
- DZone articles on Kubernetes: <https://dzone.com/users/4519738/rajeshg007.html>
- Explanation with code: <https://github.com/brainupgrade-in/kubernetes>
- Exam Syllabus: <https://github.com/cncf/curriculum>



Recap

- Review
- Q & A

Thank You for your active participation!

info@brainupgrade.in

<https://www.facebook.com/brainupgrade.in>

<https://www.linkedin.com/company/brains-upgrade/>

<https://www.linkedin.com/in/rajesh-ghaware/>
